

```

In [1]: # 1. Perform necessary data preprocessing, e.g. removing punctuation and stop words
# You may use the outputs from previous weekly assignments.
# Cleaning the input data file

#removing all tags from the input file
import re
#using nltk
import nltk
from nltk.stem.porter import PorterStemmer
from nltk import word_tokenize
import pickle
from nltk.corpus import stopwords
nltk_stopwords = nltk.corpus.stopwords.words('english')
from nltk.stem.wordnet import WordNetLemmatizer
from nltk import word_tokenize, pos_tag
from nltk.corpus import wordnet
from collections import defaultdict
import re
from nltk.stem import WordNetLemmatizer
from spellchecker import SpellChecker

# Read the input text file
with open('E:\\MS_Course_Notes\\COMP_293C\\Assignments\\Assignment_6\\webpage_data.
    text = file.read()

# Define regular expressions to match the JavaScript functions to remove
functions_to_remove = [
    r'Search{"path":.*?}\\Skip to main', # Match and remove JavaScript variables
    r'{"path":.*?}\\Skip to main', # Match and remove JavaScript variables
    r'function focusIt\\(\\) \\{\\[\\s\\S]*?\\}\\}\\);Skip to main',
    r'function focusIt\\(\\) \\{\\[\\s\\S]*? \\}\\}\\);',
    r'\\$\\(.*?\\}\\}\\);',
    r'window\\.dataLayer = window\\.dataLayer \\|\\| \\[\\];\\s*function gtag\\(\\)\\{dataLay
]

# Remove the matched functions
for pattern in functions_to_remove:
    text = re.sub(pattern, '', text, flags=re.DOTALL)

tokens_text = word_tokenize(text)

#print(tokens_text)

# Remove non-alphanumeric items, empty strings, and whitespace strings
filtered_list = [item for item in tokens_text if re.match(r'^[a-zA-Z0-9]*$', item)]

# Remove stopwords
stop_words = set(stopwords.words('english'))
filtered_list = [word for word in filtered_list if word not in stop_words]
filtered_list = [word.lower() for word in filtered_list]

# Print the filtered list without stopwords
#print(filtered_list)

lemmatizer = WordNetLemmatizer()

# Perform Lemmatization
lemmatized_list = [lemmatizer.lemmatize(word) for word in filtered_list]

# Print the stemmed and Lemmatized list
#print(lemmatized_list)

```

```
#1. Cleaning the input file without punctuation and stop words in the data files.

# Specify the path for the output file to save the tokenized text
output_file_path = "E:\\MS_Course_Notes\\COMP_293C\\Assignments\\Assignment_6\\q1_c

# Write the tokenized text to the output file
with open(output_file_path, 'w') as output_file:
    output_file.write(str(lemmatized_list))

print("Remove punctuation and stop words in the data files text written to:", output_file_path)

Remove punctuation and stop words in the data files text written to: E:\\MS_Course_Notes\\COMP_293C\\Assignments\\Assignment_6\\q1_cleaned_output.txt
```

2. Propose a binary classification problem from your project data and identify the columns that you will use to solve the problem. You may need to create new columns of data.

Binary Classification Problem: Identifying Programs and general University data

Classification is based on the Programs

As project is a university chatbot I am trying to create a binary classification model to determine whether the text refers to a program related data or general university data. The data set has the general data like pacific card, housing, dinning and course related information. So here I splitting the academic related data and non academic data in separate labels.

Dataset

The a dataset that contains various paragraphs related to educational programs and general FAQ in the university. Each paragraph is associated with a label indicating whether it's about a "non academic" (Class 0) or an "academic" (Class 1).

Columns Used

1. **Text Data (Paragraphs):** We use the text data column containing the paragraphs that describe different educational programs. This text data is essential for extracting information to make classification decisions.
2. **Target Label (Class):** We have a target label column (Class) to identify whether each paragraph corresponds to a "academic" or an "non academic."

Explanation

After calculating the frequency of academic keywords and non academic keywords so I am proceeding with those words. So that the dataset is balanced.

In [2]: *#calculating with the frequency of the words to proceed with the binary classifica*

```
import re
import nltk
from nltk.corpus import stopwords
from collections import Counter

# Tokenize the text into words
words = nltk.word_tokenize(text)

# Convert words to lowercase
words = [word.lower() for word in words]

# Remove stopwords and punctuation
stop_words = set(stopwords.words("english"))
words = [word for word in lemmatized_list if word.isalpha() and word not in stop_wc

# Count word frequencies
word_counts = Counter(words)

# Find the most common words
most_common_words = word_counts.most_common(10) # Change the number as needed

# Print the most common words
for word, count in most_common_words:
    print(f"{word}: {count}")
```

```
student: 878
program: 569
pacific: 348
application: 341
university: 287
graduate: 277
course: 234
admission: 229
scholarship: 221
eligible: 216
```

In [3]: *# Read the input text file*

```
with open('E:\MS_Course_Notes\COMP_293C\Assignments\Project\Full_Cleaned_input
    text = file.read()

input_text = text

# Add spaces after question marks and appropriate punctuation at the end of each li
formatted_text = re.sub(r'(\w)\?(\w)', r'\1? \2', input_text)
formatted_text = re.sub(r'(\w)\.', r'\1. ', formatted_text)

# Print the updated input data
#print(formatted_text)
```

In [4]: *#creating a new column for the binary classification*

```
import re

# Define the keywords related to academic
keywords = ['graduate', 'undergraduate', 'program', 'degree', 'academic',
            'student', 'application', 'international', 'scholarship', 'eligible',
            'admission', 'financial', 'aid', 'research', 'course', 'school',
            'Bachelor', 'master', 'academic', 'semester', 'eligibility', 'science']
```

```

        'education','prerequisite','gpa','research','score','fee']

# Define a function to check if any of the keywords are in a paragraph
def contains_keyword(paragraph):
    for keyword in keywords:
        if re.search(keyword, paragraph, re.IGNORECASE):
            return 1 # Label as 1 if a keyword is found
    return 0 # Return None for paragraphs without either keyword

# Split the text data into paragraphs (assuming paragraphs are separated by a period)
paragraphs = re.split(r'(?<=[.!?])\s+', formatted_text)

# Create a list of (label, paragraph) pairs, excluding None values
labeled_paragraphs = [(label, paragraph) for label, paragraph in [(contains_keyword(keyword) for keyword in keywords) for paragraph in paragraphs]]

# Print the labeled paragraphs
print(labeled_paragraphs)

# Initialize counters for labels
count_0 = 0
count_1 = 0

# Iterate through the labeled paragraphs to count the labels
# showing the dataset with is almost balanced
for label, _ in labeled_paragraphs:
    if label == 0:
        count_0 += 1
    elif label == 1:
        count_1 += 1

# Print the counts
print("Count of label 0:", count_0)
print("Count of label 1:", count_1)

```

Count of label 0: 4362
Count of label 1: 4495

```

In [5]: # Write the Labeled data to a new file
with open("E:\\MS_Course_Notes\\COMP_293C\\Assignments\\Assignment_6\\labeled_data.txt", "w") as file:
    for label, paragraph in labeled_paragraphs:
        file.write(f"{label}\t{paragraph}\n")

```

```

In [6]: #1. Splitting the data set into test and train dataset
from sklearn.model_selection import train_test_split

# Split the dataset into features (X) and labels (y)
X = [paragraph for label, paragraph in labeled_paragraphs]
y = [label for label, paragraph in labeled_paragraphs]

# Split the dataset into training and test sets (adjust the test_size as needed)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

In [7]: #3. Compute TF-IDF vectors on the text data.
from sklearn.feature_extraction.text import TfidfVectorizer

# Define the TF-IDF vectorizer
tfidf_vectorizer = TfidfVectorizer(stop_words='english')

# Fit and transform the vectorizer on the training data
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)

# Transform the test data using the same vectorizer
X_test_tfidf = tfidf_vectorizer.transform(X_test)

```

In [8]: *#4. Solve your binary classification problem with the Naïve Bayes classifier.*

```
from sklearn.naive_bayes import MultinomialNB

# Create a Naïve Bayes classifier
naive_bayes_classifier = MultinomialNB()

# Train the classifier using the TF-IDF vectors and corresponding labels
naive_bayes_classifier.fit(X_train_tfidf, y_train)

# Make predictions on the test data
y_pred = naive_bayes_classifier.predict(X_test_tfidf)
```

In [9]:

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Predict using the trained Naïve Bayes classifier and the TF-IDF vectors for test
nb_y_pred = naive_bayes_classifier.predict(X_test_tfidf)

# Calculate accuracy for the Naïve Bayes classifier
nb_accuracy = accuracy_score(y_test, nb_y_pred)
print("Naïve Bayes Accuracy:", nb_accuracy)

# Generate a classification report for the Naïve Bayes classifier
nb_report = classification_report(y_test, nb_y_pred)
print("Naïve Bayes Classification Report:\n", nb_report)


# Generate the confusion matrix for the Naïve Bayes classifier
nb_cm = confusion_matrix(y_test, nb_y_pred)

# Plot the confusion matrix for the Naïve Bayes classifier
plt.figure(figsize=(8, 6))
sns.heatmap(nb_cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Non Academic",
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix for Naïve Bayes Classifier")
plt.show()
```

Naïve Bayes Accuracy: 0.7951467268623025

Naïve Bayes Classification Report:

	precision	recall	f1-score	support
0	0.93	0.63	0.75	873
1	0.73	0.95	0.82	899
accuracy			0.80	1772
macro avg	0.83	0.79	0.79	1772
weighted avg	0.83	0.80	0.79	1772



Confusion Matrix Report Explanation

- The model correctly identified 862 cases of "non academic" (True Positives).
- The model correctly identified 513 cases of "academic" (True Negatives).
- The model incorrectly identified 360 cases as "non academic" when they were actually "academic" (False Positives).
- The model incorrectly identified 37 cases as "academic" when they were actually "non academic" (False Negatives).

Accuracy The accuracy predicted from the model is 79.5%.

```
In [10]: #5. Solve your binary classification problem with the SVC classifier.
from sklearn.svm import SVC

# Initialize the SVC classifier
svc_classifier = SVC(kernel='linear', C=1.0)

# Fit the classifier to your training data with TF-IDF vectors
svc_classifier.fit(X_train_tfidf, y_train)
```

```
Out[10]: SVC
SVC(kernel='linear')
```

```
In [11]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Predict using the trained classifier and the TF-IDF vectors for test data
y_pred = svc_classifier.predict(X_test_tfidf)
```

```
# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a classification report
report = classification_report(y_test, y_pred)
print("Classification Report:\n", report)


# Generate the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix for the SVM classifier
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Non Academic", "Academic"], yticklabels=["Non Academic", "Academic"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix for SVM Classifier")
plt.show()
```

Accuracy: 0.9644469525959368

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.99	0.96	873
1	0.99	0.94	0.96	899
accuracy			0.96	1772
macro avg	0.97	0.96	0.96	1772
weighted avg	0.97	0.96	0.96	1772



Confusion Matrix Report Explanation

- The model correctly identified 863 cases as "academic" (True Positives).
- The model correctly identified 846 cases as "non academic" (True Negatives).
- The model incorrectly identified 10 cases as "academic" when they were actually "non academic" (False Positives).
- The model incorrectly identified 53 cases as "non academic" when they were actually "academic" (False Negatives).

Accuracy The accuracy predicted from the model is 96.4%.